

Week 02, Monday

Memory and Data

The C View of Data

2/30

A C program sees data as a collection of *variables*

```
int g = 2;

int main(void)
{
    int i;
    int v[5]
    char *s = "Hello";
    int *n = malloc(sizeof(int));
    ...
}
```

Each variable has a number of properties (e.g. name, type, size)

... The C View of Data

3/30

Variables are examples of *computational objects*

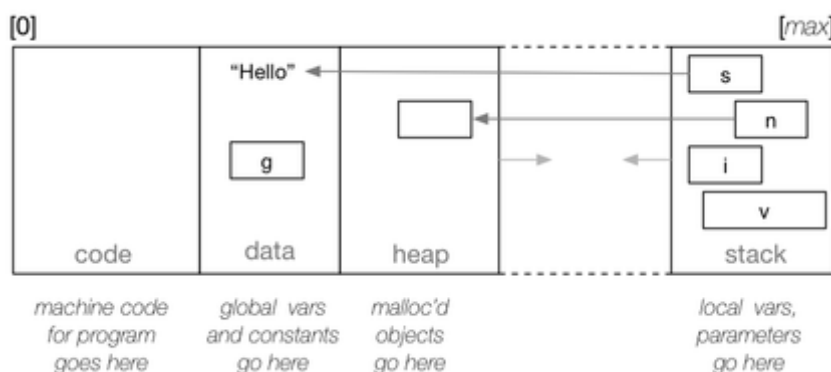
Each computational object has

- a *location* in memory
- a *value* (ultimately just a bit-string)
- a *name* (unless created by `malloc()`)
- a *type*, which determines ...
 - its *size* (in units of whole bytes, `sizeof`)
 - how to *interpret* its value
 - what *operations* apply to the value
- a *scope* (where it's visible within the program)
- a *lifetime* (during which part of program execution it exists)

... The C View of Data

4/30

C allocates data objects to various well-defined regions of memory during program execution



Exercise 1: Properties of Variables

5/30

Identify the properties of each of the named objects in the following:

```

int a;           // global int variable

int main(void) {
    int b;       // local int variable
    char c;      // local char variable
    char d[10];  // local char array
    ...
}

int e;           // global? int variable


int f(int g) {   // function + parameter
    double h;    // local double variable
    ...
}

```

The Physical View of Data

6/30

Memory = indexed array of bytes



Indexes are "memory addresses" (a.k.a. pointers)

Data can be fetched in chunks of 1,2,4,8 bytes

Memory

7/30

Also called: RAM, main memory, primary storage, ...

Technology: semiconductor-based

Distinguishing features

- relatively large (e.g. 2^{28} bytes)
- any byte can be fetched with same cost
- cost of fetching 1,2,4,8 bytes is small (ns)

Two properties related to data persistence

- *volatile* (e.g. DRAM) ... data lost when powered off
- *non-volatile* (e.g. EEPROM) ... data stays when powered off

... Memory


8/30

When addressing objects in memory

- any byte address can be used to fetch 1-byte object
- byte address for N -byte object must be divisible by N

Data is fetched into N -byte CPU registers for use

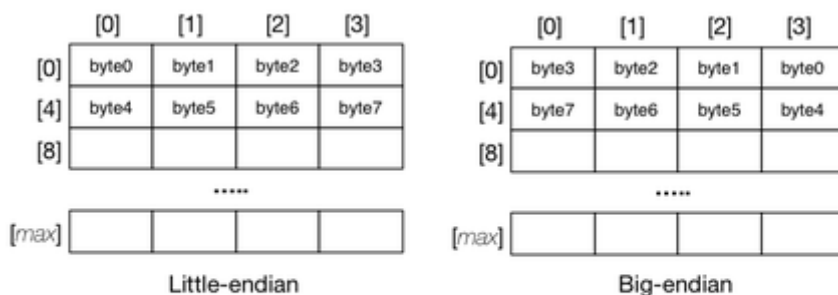
Data bytes in registers may be in different order to memory, e.g.




... Memory

9/30

Memories can be categorised as *big-endian* or *little-endian*



Loading a 4-byte int from address 0 gives



Data Representation

Data Representation

11/30

Ultimately, memory allows you to

- load bit-strings of sizes 1,2,4,8 bytes
- from N -byte boundary addresses
- into registers in the CPU

What you are presented with is a string of 8,16,32,64 bits

Need to *interpret* this bit-string as a meaningful value

Data representations provide a way of assigning meaning to bit-strings

Character Data

12/30

Character data has several possible representations (encodings)

The two most common:

- ASCII (ISO 646)
 - 7-bit values, using lower 7-bits of a byte (top bit always zero)
 - can encode roman alphabet, digits, punctuation, control chars
- UTF-8 (Unicode)
 - 8-bit values, with ability to extend to multi-byte values
 - can encode all human languages plus other symbols

(e.g. $\sqrt{\sum \forall \exists}$ or      )

ASCII Character Encoding

Uses values in the range 0x00 to 0x7F (0..127)

Characters partitioned into sequential groups

- control characters (0..31) ... e.g. '\0', '\n'
- punctuation chars (32..47, 91..96, 123..126)
- digits (48..57) ... '0'..'9'
- upper case alphabetic (65..90) ... 'A'..'Z'
- lower case alphabetic (97..122) ... 'a'..'z'

In C, can map between char and ascii code by e.g. ((int)'a')

Sequential nature of groups allow for e.g. (ch - '0')

... ASCII Character Encoding

14/30

Hexademical ASCII char table (from `man 7 ascii`)

| | | | | | | | |
|--------|--------|--------|--------|--------|--------|--------|--------|
| 00 nul | 01 soh | 02 stx | 03 etx | 04 eot | 05 enq | 06 ack | 07 bel |
| 08 bs | 09 ht | 0a nl | 0b vt | 0c np | 0d cr | 0e so | 0f si |
| 10 dle | 11 dc1 | 12 dc2 | 13 dc3 | 14 dc4 | 15 nak | 16 syn | 17 etb |
| 18 can | 19 em | 1a sub | 1b esc | 1c fs | 1d gs | 1e rs | 1f us |
| 20 sp | 21 ! | 22 " | 23 # | 24 \$ | 25 % | 26 & | 27 ' |
| 28 (| 29) | 2a * | 2b + | 2c , | 2d - | 2e . | 2f / |
| 30 0 | 31 1 | 32 2 | 33 3 | 34 4 | 35 5 | 36 6 | 37 7 |
| 38 8 | 39 9 | 3a : | 3b ; | 3c < | 3d = | 3e > | 3f ? |
| 40 @ | 41 A | 42 B | 43 C | 44 D | 45 E | 46 F | 47 G |
| 48 H | 49 I | 4a J | 4b K | 4c L | 4d M | 4e N | 4f O |
| 50 P | 51 Q | 52 R | 53 S | 54 T | 55 U | 56 V | 57 W |
| 58 X | 59 Y | 5a Z | 5b [| 5c \ | 5d] | 5e ^ | 5f _ |
| 60 ` | 61 a | 62 b | 63 c | 64 d | 65 e | 66 f | 67 g |
| 68 h | 69 i | 6a j | 6b k | 6c l | 6d m | 6e n | 6f o |
| 70 p | 71 q | 72 r | 73 s | 74 t | 75 u | 76 v | 77 w |
| 78 x | 79 y | 7a z | 7b { | 7c | 7d } | 7e ~ | 7f del |

0x0a = '\n', 0x20 = ' ', 0x09 = '\t', but note no EOF

Exercise 2: Using 'a'..'z' as indexes

15/30

Write C code that allows you to treat an array like

```
int freq[26];
```

as if it were indexed by 'a'..'z'

Sample usage

```
for (char c = 'a'; c <= 'z'; c++)
    freq[XXX] = 0;
...
for (char c = 'a'; c <= 'z'; c++)
    printf("%s has freq %d\n", c, freq[XXX]);
```

In other words, replace the xxx by an index calculation

UTF-8 Character Encoding

16/30

UTF-8 uses a variable-length encoding as follows

| #bytes | #bits | Byte 1 | Byte 2 | Byte 3 | Byte 4 |
|--------|-------|----------|--------|--------|--------|
| 1 | 7 | 0xxxxxxx | - | - | - |

| | | | | | |
|---|----|----------|----------|----------|----------|
| 2 | 11 | 110xxxxx | 10xxxxxx | - | - |
| 3 | 16 | 1110xxxx | 10xxxxxx | 10xxxxxx | - |
| 4 | 21 | 11110xxx | 10xxxxxx | 10xxxxxx | 10xxxxxx |

The 127 1-byte codes are compatible with ASCII

The 2048 2-byte codes include most Latin-script alphabets

The 65536 3-byte codes include most Asian languages

The 2097152 4-byte codes include symbols and emojis and ...

... UTF-8 Character Encoding

17/30

UTF-8 examples

| ch | unicode | bits | simple binary | UTF-8 binary |
|----|---------|------|----------------------------|-------------------------------------|
| \$ | U+0024 | 7 | 010 0100 | 00100100 |
| ¢ | U+00A2 | 11 | 000 1010 0010 | 11000010 10100010 |
| € | U+20AC | 16 | 0010 0000 1010 1100 | 11100010 10000010 10101100 |
| ☐ | U+10348 | 21 | 0 0001 0000 0011 0100 1000 | 11110000 10010000 10001101 10001000 |

Unicode strings can be manipulated in C (e.g. "안녕하세요")

Like other C strings, they are terminated by a 0 byte (i.e. '\0')

... UTF-8 Character Encoding

18/30

Unicode constants in C strings ...

The following two notations work in some contexts

- `\uHexDigits` ... insert Unicode code value
- `\x2HexDigits` ... insert individual bytes

Examples:

```
"abc\xE2\x86\xABdef"
u"abc\u21ABdef"
```


The red sequences produce 3 bytes and 1 Unicode symbol.

Numeric Data

19/30

Numeric data comes in two major forms

- integer ... subset (range) of the mathematical integers
- floating point ... subset of the mathematical real numbers



20/30

Integer Constants

Three ways to write integer constants in C

- 42 ... signed decimal (0..9)
- 0x2A ... unsigned hexadecimal (0..F)
- 052 ... signed octal (0..7)

Variations

- 123U ... unsigned int value (typically 32 bits)
- 123L ... long int value (typically 64 bits)
- 123S ... short int value (typically 16 bits)

Invalid constants lie outside the range for their type, e.g.


- 4294967296, -1U, 6666666S, 078

Unsigned integers

21/30

The **unsigned int** data type

- commonly 32 bits, storing values in the range $0 \dots 2^{32}-1$



... Unsigned integers

22/30

Value interpreted as binary number

E.g. consider an 8-bit unsigned int

$$01001101 = 2^6 + 2^3 + 2^2 + 2^0 = 64 + 8 + 4 + 1 = 77$$

Addition is bitwise with carry

| | | | |
|------------|------------|------------|------------|
| 00000001 | 00000001 | 01001101 | 11111111 |
| + 00000010 | + 00000011 | + 00001011 | + 00000001 |
| ----- | ----- | ----- | ----- |
| 00000011 | 00000100 | 01011000 | 00000000 |

Most machines will also flag the *overflow* in the fourth example

Exercise 3: Binary↔decimal Conversion

23/30

Convert these 8-bit binary numbers to hexadecimal:

- 00001001, 00001101, 00101010, 00110011, 11001100

Convert these 8-bit binary numbers to decimal:

- 00001001, 00001101, 00101010, 00110011, 11001100

Convert the following decimal numbers to 8-bit binary:


- 15, 64, 99, 200, 256

Signed integers

24/30

The **int** data type

- commonly 32 bits, storing values in the range $-2^{31} \dots 2^{31}-1$



... Signed integers

25/30

Several possible representations for negative values

- signed magnitude ... first bit is sign, rest are magnitude
- ones complement ... form $-N$ by inverting all bits in N
- twos complement ... form $-N$ by inverting N and adding 1

In all representations, +ve numbers have 0 in leftmost bit

Examples: representations of (8-bit) -5 (where 5 is 00000101)

- 10000101 ... signed magnitude
- 11111010 ... ones complement
- 11111011 ... twos complement

... Signed integers

26/30

Signed magnitude: Easy to form $-X$ from X ... OR in high-order bit

A problem (using 8-bit ints) ...

- what do these numbers represent? 00000000, 10000000

Two zeroes ... one positive, one negative

Another problem: $x + -x \neq 0$ (mostly) with simple addition

| | | | | | |
|------------|----|------------|-----|------------|------|
| 00000011 | 3 | 00101010 | 42 | 01111111 | 127 |
| + 10000011 | -3 | + 10101010 | -42 | + 11111111 | -127 |
| ----- | | ----- | | ----- | |
| 10000110 | !0 | 11010100 | !0 | 01111110 | !0 |

To fix requires extra hardware in ALU

... Signed integers

27/30

Ones complement: Easy to form $-X$ from X ... NEG all bits

A problem (using 8-bit ints) ...

- what do these numbers represent? 00000000, 11111111

Two zeroes ... one positive, one negative

At least $x + -x$ is equal to one of the zeroes with simple addition

| | | | | | |
|------------|----|------------|-----|------------|----|
| 00000011 | 3 | 00101010 | 42 | 01111111 | |
| + 11111100 | -3 | + 11010101 | -42 | + 10000000 | |
| ----- | | ----- | | ----- | |
| 11111111 | !0 | 11111111 | !0 | 11111111 | -0 |

... Signed integers

28/30

Twos complement: to form $-X$ from X ... NEG all bits, then add 1

Now have only one representation for zero (00000000)

$$\circ -0 = \sim 00000000 + 1 = 11111111 + 1 = 00000000$$

Only one zero value. Also, $-(-x) = x$

Even better, $x + -x = 0$ in all cases with simple addition

| | | | | | |
|------------|----|------------|-----|------------|---|
| 00000011 | 3 | 00101010 | 42 | 01111111 | |
| + 11111101 | -3 | + 11010110 | -42 | + 10000001 | |
| ----- | | ----- | | ----- | |
| 00000000 | 0 | 00000000 | 0 | 00000000 | 0 |

Always produces an "overflow" bit, but can ignore this

Exercise 4: Binary↔decimal Conversion

29/30

What decimal numbers do these 8-bit twos complement numbers represent:

$$\circ 10001001, 10001101, 10101010, 10110011, 11001100$$

Convert the following decimal numbers to 8-bit binary:

$$\circ 15, 64, 99, 127, 128$$

Show signed magnitude, 1's complement and 2's complement

Demonstrate the addition of $x + -x$, where x is

$$\circ 5, 20, 64, 99, 127$$

Exercise 5: Integer Powers

30/30

C does not have a power operator (e.g. like $x**y = x^y$)

Write a function to compute x^y

```
int raise(int x, int y) { ... }
```

Write a specialised version to compute 2^y

```
int powOf2(int y) { ... }
```

Produced: 6 Aug 2017