

# Week 03

## Data Representation (cont)

### The story so far ...

2/36

We have looked at:

- character data
  - ASCII ... 7-bit encoding of English alphabet
  - UTF8 ... 7,11,16,21 bit encoding of all languages
- integer data
  - unsigned int ... 32-bit encoding of  $0..2^{32}-1$
  - int ... 32-bit twos-complement encoding of  $-2^{31}..2^{31}-1$

### Pointers

3/36

Pointers represent memory addresses/locations

- number of bits depends on memory size, but typically 32-bits
- data pointers reference addresses in *data/heap/stack* regions
- function pointers reference addresses in *code* region

Many kinds of pointers, one for each data type, but

- `sizeof(int *) = sizeof(char *)`  
`= sizeof(double *) = sizeof(struct X *)`

Pointer *values* must be appropriate for data type, e.g.

- `(char *)` ... can reference any byte address
- `(int *)` ... must have `addr % 4 == 0`
- `(double *)` ... must have `addr % 8 == 0`

### Exercise 1: Valid Pointers

4/36

Which of the following are likely to be valid pointers

```
0x00000000    0x00001000    0x00001001
0x7f000000    0x7f000001    0x7f000004
```

to objects of type

- char
- int
- unsigned int
- double
- int (\*f)()

### ... Pointers

5/36

Can "move" from object to object by *pointer arithmetic*

For any pointer `T *p`, `p++` increases `p` by `sizeof(T)`

Examples (assuming 16-bit pointers):

```
char *p = 0x6060; p++; assert(p == 0x6061)
int *q = 0x6060; q++; assert(q == 0x6064)
double *r = 0x6060; r++; assert(r == 0x6068)
```

A common (efficient) paradigm for scanning a string

```
char *s = "a string";
char *c;
// print a string, char-by-char
for (c = s; *c != '\0'; c++) {
    printf("%c", *c);
}
```

## Exercise 2: Sum an array of ints

6/36

Write a function

```
int sumOf(int *a, int n) { ... }
```

to sum the elements of array `a[ ]` containing `n` values.

Implement it two ways:

- using the "standard" approach with an index
- using a pointer that scans the elements

## Floating Point Numbers


7/36

Floating point numbers model a (tiny) subset of  $\mathbb{R}$

- many real values don't have exact representation (e.g.  $1/3$ )
- results of calculations may contain small inaccuracies

*Precision* categorises how close to exact

- numbers close to zero have higher precision (more accurate)
- numbers further from zero have lower precision (less accurate)



### ... Floating Point Numbers

8/36

C has two floating point types

- **float** ... typically 32-bit quantity (lower precision, narrower range)
- **double** ... typically 64-bit quantity (higher precision, wider range)

Literal floating point values: `3.14159`, `1.0/3`, `1.0e-9`

Display via `printf`

```
printf("%w.Pf", (float)2.17828)
printf("%w.PlF", (double)2.17828)
```

`w` gives total width (blank padded), `P` gives #digits after dec point

```
printf("%10.4lf", (double)2.718281828459);
displays  _2.7183
printf("%20.20lf", (double)4.0/7);
displays 0.57142857142857139685
```

### ... Floating Point Numbers

9/36

IEEE 754 standard ...

- scientific notation with *fraction*  $F$  and *exponent*  $E$
- numbers have form  $F \times 2^E$ , where both  $F$  and  $E$  can be -ve

- INFINITY = representation for  $\infty$  and  $-\infty$  (e.g. 1.0/0)
- NAN = representation for invalid value NaN (e.g. sqrt(-1.0))
- 32-bit single-precision, 64-bit double precision

Fraction part is *normalised* (i.e.  $1.2345 \times 10^2$  rather than 123.45)


Example of normalising in binary:

- 1010.1011 is normalized as  $1.0101011 \times 2^{011}$
- $1010.1011 = 10 + 11/16 = 10.6875$
- $1.0101011 \times 2^{011} = (1 + 43/128) * 2^3 = 1.3359375 * 8 = 10.6875$

### ... Floating Point Numbers

10/36

Internal structure of floating point values



More complex than int because *1.dddd e dd*

### ... Floating Point Numbers

11/36

Details of internal structure

- fraction part is always *1.bbbbbbbb*; don't store 1
- exponent is offset relative to a baseline  $-2^{b-1}-1$ ,  $b$  = #exponent bits

Ranges of values for 32-bit single-precision float:

Component	Min Value	Max Value
exponent	00000000 = -127	11111111 = 128
fraction	00...00 = 0	11...11 = $2^{-1} + 2^{-2} + \dots + 2^{-24}$

00...00 = 24 zero bits, 11...11 = 24 one bits

Ranges of values for 64-bit double-precision float:

Component	Min Value	Max Value
exponent	000000000000 = -2047	111111111111 = 2048
fraction	00...00 = 0	11...11 = $2^{-1} + 2^{-2} + \dots + 2^{-51}$

00...00 = 24 zero bits, 11...11 = 24 one bits

### ... Floating Point Numbers

12/36

Example (single-precision):

```

150.75 = 10010110.11
        // normalise fraction, compute exponent
        = 1.001011011 × 27
        // determine sign bit,
        // map fraction to 24 bits,
        // map exponent relative to baseline
        = 01000001100010110110000000000000

```

where red is sign bit, green is exponent, blue is fraction

Note:

- the baseline (aka bias) is 127, the exponent is 2<sup>7</sup>
- so, in the exponent, we store 127+7 = 134 = 10000110

### Exercise 3: Floating point → Decimal

13/36

Convert the following floating point numbers to decimal.

Assume that they are in IEEE 754 single-precision format.

0 10000000 110000000000000000000000

1 01111110 100000000000000000000000

### Arrays

14/36

Arrays are defined to have  $N$  elements, each of type  $T$


Examples:

```

int    a[100];    // array of 10 ints
char   str[256];  // array of 256 chars
double vec[100];  // array of 100 doubles

```

Elements are laid out adjacent in memory



### ... Arrays

15/36

Assuming an array declaration like `Type v[N]` ...

- individual array elements are accessed via indices  $0..N-1$
- total amount of space allocated to array  $N \times \text{sizeof}(Type)$

Name of array gives address of first element (e.g. `v = &v[0]`)

Name of array can be treated as a pointer to element type `Type`

Array indexing can be treated as `v[i] ≡ *(v+i)`

If have pointer to first element, can use it just like an array


Strings are just arrays of `char` with a `'\0'` terminator

- constant strings have `'\0'` added automatically
- string buffers must allow for element to hold `'\0'`

## ... Arrays

16/36

When arrays are "passed" to a function, actually pass `&a[0]`



## Exercise 4: Initialising Strings

17/36

Explain the difference between the following initialisers:

```

char a[9] = "a string";

char *b = "a string";

char c[9] = {'a', ' ', 's', 't', 'r', 'i', 'n', 'g', '\0'};

char *d = {'a', ' ', 's', 't', 'r', 'i', 'n', 'g', '\0'};

```

## ... Arrays

18/36

Arrays can be created automatically or via `malloc()`

```

int main(void)
{
    char str1[9] = "a string";
    char *str2; // no array object yet

    str2 = malloc(20*sizeof(char));
    strcpy(str2, str);
    printf("&str1=%p, str1=%s\n", str1, str1);
    printf("&str2=%p, str2=%s\n", str2, str2);

    free(str2);
    return 0;
}

```

Two separate arrays (different `&s`), but have same contents

(except for the uninitialised parts of the arrays)

## Structs

19/36

Structs are defined to have a number of components


- each component has a *Name* and a *Type*

Example:

```
typedef struct ... Date;

struct _student {
    int id;
    char given[50];
    char family[50];
    char gender;
    Date bday;
};
```

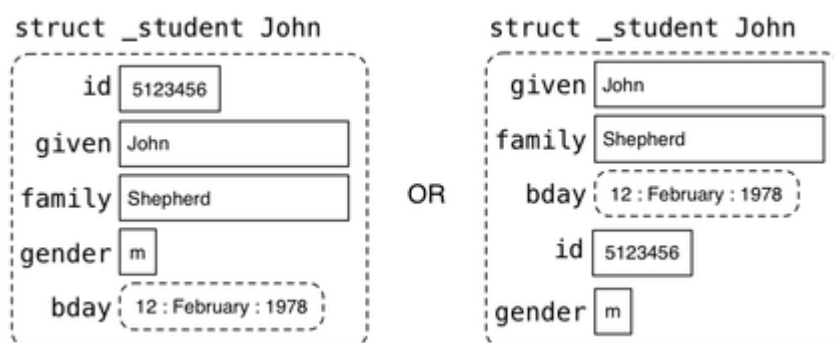
*Defines a new data type called struct \_student*



## ... Structs

20/36

Internal layout of struct components determined by compiler




Each name maps to a byte offset within the struct

E.g. in first example id = offset 0, given = offset 4, family = offset 54, etc.

## ... Structs

21/36

To ensure alignment, internal "padding" may be needed



Padding wastes space; re-order fields to minimise waste.

clang has `-Wpadded` to warn about padding in structs

## Exercise 5: Struct Internal Layout

22/36

Determine the offsets of the fields in


```
struct _s1 {
    char a[6]; // array of 6 1-byte chars
    int b;    // 4-byte int
    char c;    // 1-byte char
```

```
double d;    // 8-byte int
int e;       // 4-byte int
char f;      // 1-byte char
};
```

## Exercise 6: Struct Alternatives

23/36

Consider these two possible representations of Nodes in a linked list of strings:



Show how they would be (a) defined, (b) initialised. How large is each?

## Variable-length Structs

24/36

Structs can contain pointers to dynamic objects

But we can also "embed" one dynamic object in a malloc'd struct

- define the dynamic object as the last component
- malloc() more space than the struct requires
- to make the final component as large as required

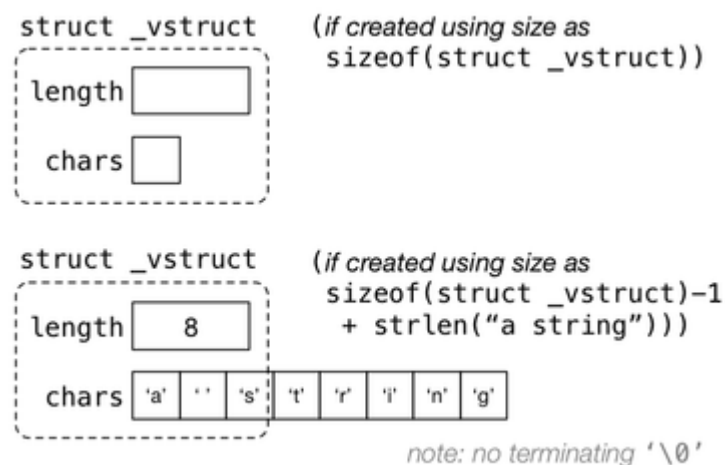
Example:

```
struct _vstruct {
    int length;
    char chars[1]; // array whose real length
                  // is calculated when malloc'ing
}
```

### ... Variable-length Structs

25/36

Amount of memory allocated to struct is determined dynamically:



## Bit-wise Structs

26/36

For fine-grained control over layout of fields in structs

- C allows programmers to specify structs bit-wise

Bit-field structs ...

- specify unnamed **components** using standard types
- specify named individual **bit fields** in each component

Example:

```
struct _bit_fields {
    unsigned int first_bit    : 1,
                  next_7_bits : 7,
                  last_24_bits : 24;
};
```

Has one component and three bit fields within that component.

### ... Bit-wise Structs

27/36

Two ways of declaring bit fields:

```
struct _bit_fields {
    unsigned int first_bit    : 1,
                  next_7_bits : 7,
                  last_24_bits : 24;
};
OR
struct _bit_fields {
    unsigned int first_bit    : 1;
    unsigned int next_7_bits  : 7;
    unsigned int last_24_bits : 24;
};
```

In both cases, `sizeof(struct _bit_fields)` is 4 bytes.

First way makes it clearer that a single `unsigned int` is used.

### ... Bit-wise Structs

28/36

Another example (graphics objects):

```
struct _object { // comprised of two 32-bit words
    unsigned int red    : 5, // 5 bits for red
                  blue   : 5, // 5 bits for blue
                  green  : 5, // 5 bits for green
                  pad    : 1, // 1 bit to pad to short
                  ident  : 16; // 16 bits for object ID
    unsigned int height : 6, // 6 bits for object height
                  width  : 6, // 6 bits for object width
                  xcoord : 9, // 9 bits for x-coordinate
                  ycoord : 9; // 9 bits for y-coordinate
};


struct _object oval;
...
oval.red = 4; oval.blue = 31; oval.green = 15;
oval.height = 5; oval.width = 15;
```

### ... Bit-wise Structs

29/36

The graphics object would be stored in memory as:





### ... Bit-wise Structs

30/36

Bit-fields provide an alternative to bit operators and masks:

```
typedef unsigned int uint;
typedef uint privs;
#define OWNER_READ (1 << 8)
#define OWNER_WRITE (1 << 7)
#define OWNER_EXEC (1 << 6)
...
#define OTHER_WRITE (1 << 1)
#define OTHER_EXEC (1 << 0)
unsigned int myPrivs;

struct _privs {
    unsigned int
        owner_read : 1,
        owner_write : 1,
        owner_exec : 1,
        ...
        other_write : 1,
        other_exec : 1;
} myPrivs;

// give owner execute permission on file
myPrivs |= OWNER_EXEC;
myPrivs.owner_exec = 1;

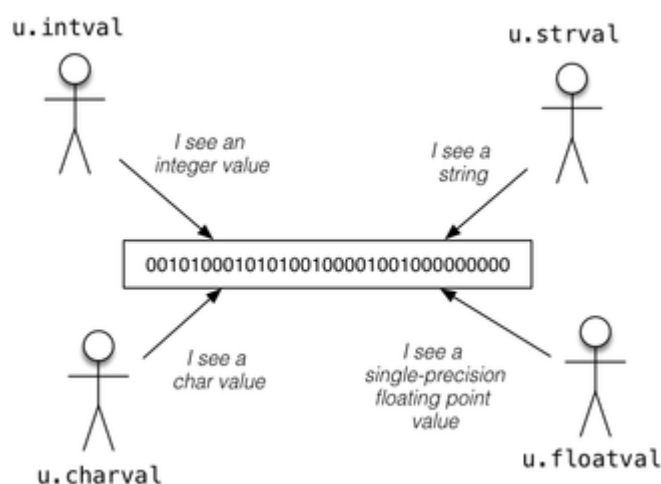
// prevent others from writing on file
myPrivs &= ~OTHER_WRITE;
myPrivs.other_write = 0;

// check whether file is readable to all
open = myPrivs & OTHER_READ;
open = myPrivs.other_read;
```

### Unions

31/36

Unions allow programmers to specify multiple interpretations for a single piece of memory.



### ... Unions

32/36

Example of defining a union type (cf. struct):

```
union _alltypes {
    int    intval;
    char  strval[4];
};
```

```
char charval;
float floatval;
};
```

```
union _alltypes myUnion;
```

myUnion is a single 4-byte memory object

Programmers can specify how to interpret bits using field names.


In the example above, all components are coincidentally the same size (4 bytes)

### ... Unions

33/36

Difference between a struct and a union

```
struct _s {          union _u {
    int    x;          int    x;
    float  y;          float  y;
} svar;              } uvar;
```



### ... Unions

34/36

General syntax for defining union types and variables:

```
union Tag {
    Type1 Member1;
    Type2 Member2;
    Type3 Member3;
    ...
} uvar;
```

*Type* can be any C type; *Member* names must be distinct

`sizeof(Union)` is the size of the largest member

`&uvar.Member1 == &uvar.Member2 == &uvar.Member3 ...`

### ... Unions

35/36

Common use of union types: "generic" variables

```
#define IS_INT    1
#define IS_FLOAT  2
#define IS_STR    3

struct _generic {
    int    vartype;
    union { int ival; char *sval; float fval; };
};
```

```
struct _generic myVar;
// treat myVar as an integer
myVar.vartype = IS_INT;
myVar.ival    = 42;
printf("%d\n", myVar.ival);
// now treat myVar as a float
```

```
myVar.vartype = IS_FLOAT;
myVar.fval    = 3.14159;
printf("%0.5f\n", myVar.fval);
```

---

## Enumerated Types

36/36

Enumerated types allow programmers to define a set of distinct named values

```
typedef enum { RED, YELLOW, BLUE } PrimaryColours;
```

```
typedef enum { LOCAL, INTL } StudentType;
```

The names are assigned consecutive `int` values, starting from 0

Above `PrimaryColors` type is equivalent to

```
#define RED    0
#define YELLOW 1
#define BLUE   2
```

Variables of type `enum...` are effectively unsigned `ints`.

---

Produced: 10 Aug 2017